
Tracking Down Software Program Bugs Using Automatic Anomaly Detection

Sudheendra Hangal, Processor Products Group, Sun

Monica S. Lam, Stanford University

ICSE-2002


~~BUENOS A~~ ORLANDO, FLORIDA

Outline

- ◆ **Background & Motivation**
- ◆ **Past Work**
- ◆ **Dynamic Invariants**
- ◆ **DIDUCE**
 - ◆ **Dynamic Invariant Detection \cup Checking Engine**
- ◆ **DIDUCE experiences**
- ◆ **Uses**
- ◆ **Conclusion, Q&A**

Background

- **Software reliability is an increasing problem**
- **Traditional input-output interaction with a program is insufficient**



- **What happens inside the black box ?**
 - **Can cheap machine cycles help ?**

Motivation for DIDUCE

- ♦ **How do you debug...**
 - ♦ **a program working correctly on some inputs, failing on others ?**
 - **"Hmm... What's different about runs which fail ?"**
 - ♦ **a program failing after a long time ?**
 - ♦ **a program you don't even know has a bug ?**
 - ♦ **Large programs written by others and evolved over time (with misleading comments!)**
- ♦ **DIDUCE successfully tackled these problems on 4 different applications we tried**

Past Work

- ♦ **Many bug detection tools**
- ♦ **Static approaches**
 - ♦ e.g. Prefix, Metal, Vault, ESC, etc
 - ♦ Exhaustive, conservative
 - ♦ No need for test inputs
- ♦ **Dynamic approaches**
 - ♦ e.g. Purify, Eraser, assert's, ...
 - ♦ Need good test input set, observe only possible behavior

Invariant Specifications

- ◆ **Many annotation based approaches**
- ◆ **Its manual work - users never do it!**
- ◆ **Usually incomplete**
- ◆ **Users may not even know the invariants...**
- ◆ **... Or may know them wrongly!**
- ◆ **Programmers rely on passing a test-suite as proof that code works**

Dynamic Invariant Detection

- ♦ Hypothesize a space of invariants associated with the program
- ♦ Test each hypothesis on program runs
 - ♦ rule out invariants which do not hold
- ♦ Prior work (Daikon) applied to small programs
- ♦ Approach is automatic and pervasive...
 - ♦ ...but may be unsound!

Dynamic Invariant Checking

- ♦ **Idea: Close the loop**
- ♦ **Detect invariants on "presumed-good" runs**
- ♦ **Automatically check invariants on other runs**
 - ♦ **Report invariant violations**
 - ♦ **Refine invariants, as you check**
- ♦ **Invariant violations signal anomalies, e.g.:**
 - ♦ **New code executed**
 - ♦ **New values seen for variables**

Dynamic Invariant Detection U Checking Engine (DIDUCE)

- ♦ Simple, practical, and effective tool
- ♦ Adds instrumentation to Java bytecode
 - ♦ To deduce and check invariants on the fly
 - ♦ Works with any compliant JVM
- ♦ 2 modes - training and checking
 - ♦ Invariant violations suppressed in training mode
 - ♦ Training continues in checking mode
- ♦ Confidence level associated with invariants and violations

GUI Screenshot

DIDUCE Invariant Violations from "buggy-run.log"

Filter Reread

```

    if (way[i] != theObj.stamps[wayAttNum]).
        throw new Exception("discharge without admit ar add");.
    head[i] += 1;.
    if (head[i] == tail[i]){.
        empty[i] = 1;.
        if (stall){.
            breadthstall = false;.
            stall = depthstall;.
        }.
    }.
    theObj.position[theObj.level] = nextPosn;.
    return nextPosn;.
}.

public int getWayNum(int slotNum){.
    if (empty[slotNum] == 0).
        return(way[slotNum]);.
    return -1;.
}.

public int getSetNum(int slotNum){.
    if (empty[slotNum] == 0).
        return(set[slotNum]);.
    return -1;.
}.

```

No.	Conf. change	What	Type	Old value	New value	Sampl...	Where	Line#
385	1 instance field read of entity...	New code	-	-	-	0	pendingBuffer.add(Lentity; I	150
386	1 array write (type integer)	New code	-	-	-	0	pendingBuffer.add(Lentity; I	150
387	-211,245.75	array read (type integer)	Value	always 0x1 (0x2	281661	pendingBuffer.discharge(Len...	164
388	-10,736,529	array read (type integer)	Value	always 0x0 (0x1	21473...	pendingBuffer.check(Lentity; I	206
389	-154,719.5	array read (type integer)	Value	always 0x0 (0x1	309439	pendingBuffer.discharge(Len...	157
390	1 param 0 of Exception.<init...	New code	-	-	-	0	pendingBuffer.discharge(Len...	158

DIDUCE Invariants

- ◆ **Values at some types of program points**
 - ◆ **Object reads/writes**
 - ◆ **Static var read/writes**
 - ◆ **Method call sites and return values**
- ◆ **Invariants associated with Tracked Expressions (TEs)**
 - ◆ **Value accessed**
 - ◆ **Change in value (for writes)**
 - ◆ **Runtime type of object being accessed**
- ◆ **Requires only "local" computation**

Tracked Expressions

Source code

```
Class SomeClass {  
  static int x;  
  int y;  
  ...  
}
```

```
Object o = new SomeClass();  
Object arr[] = new SomeClass[3];  
...
```

Tracked Expressions

Tracked Expressions

Source code

```
Class SomeClass {  
  static int x;  
  int y;  
  ...  
}
```

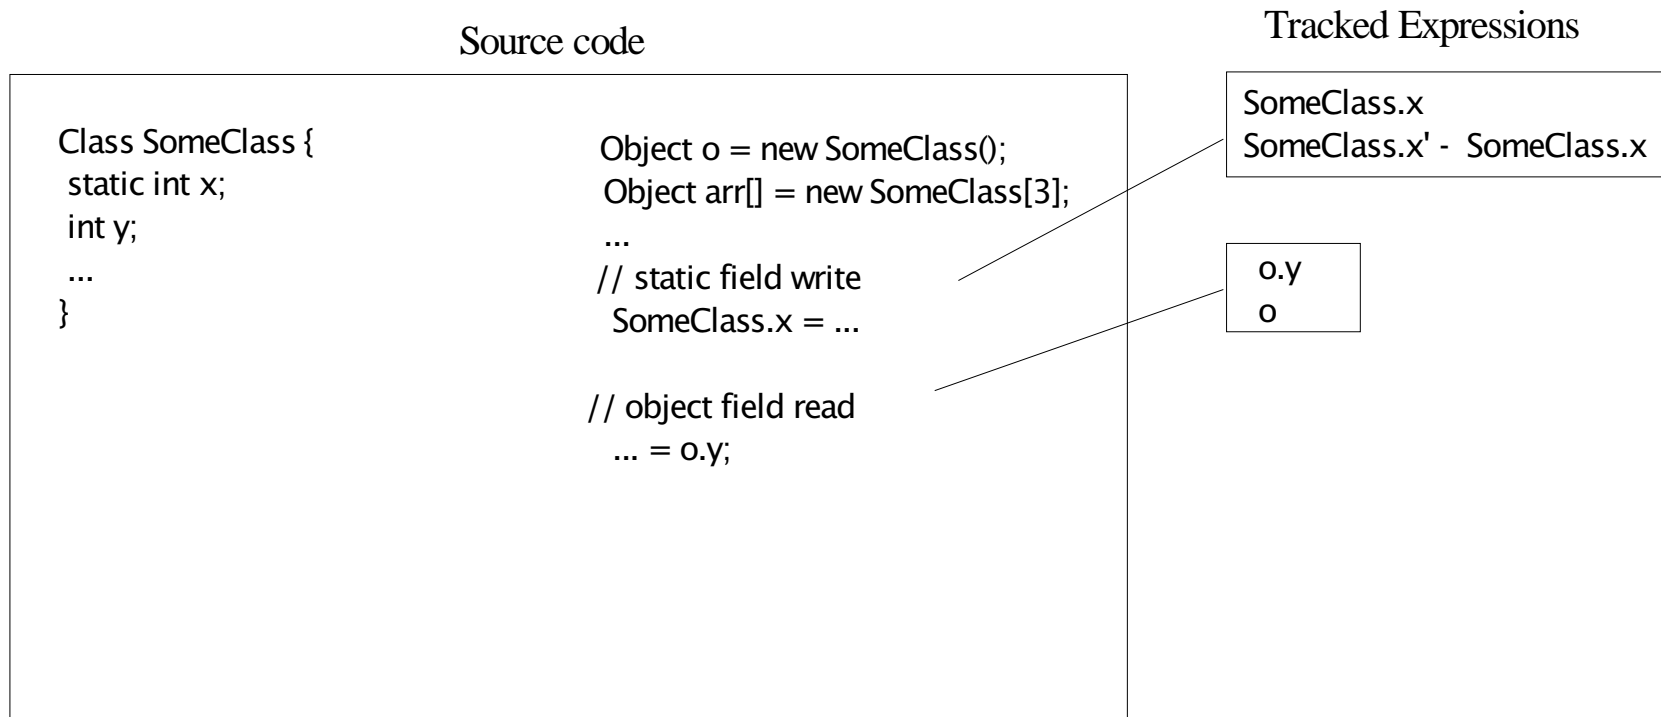
```
Object o = new SomeClass();  
Object arr[] = new SomeClass[3];  
...  
// static field write  
SomeClass.x = ...
```

Tracked Expressions

```
SomeClass.x  
SomeClass.x' - SomeClass.x
```

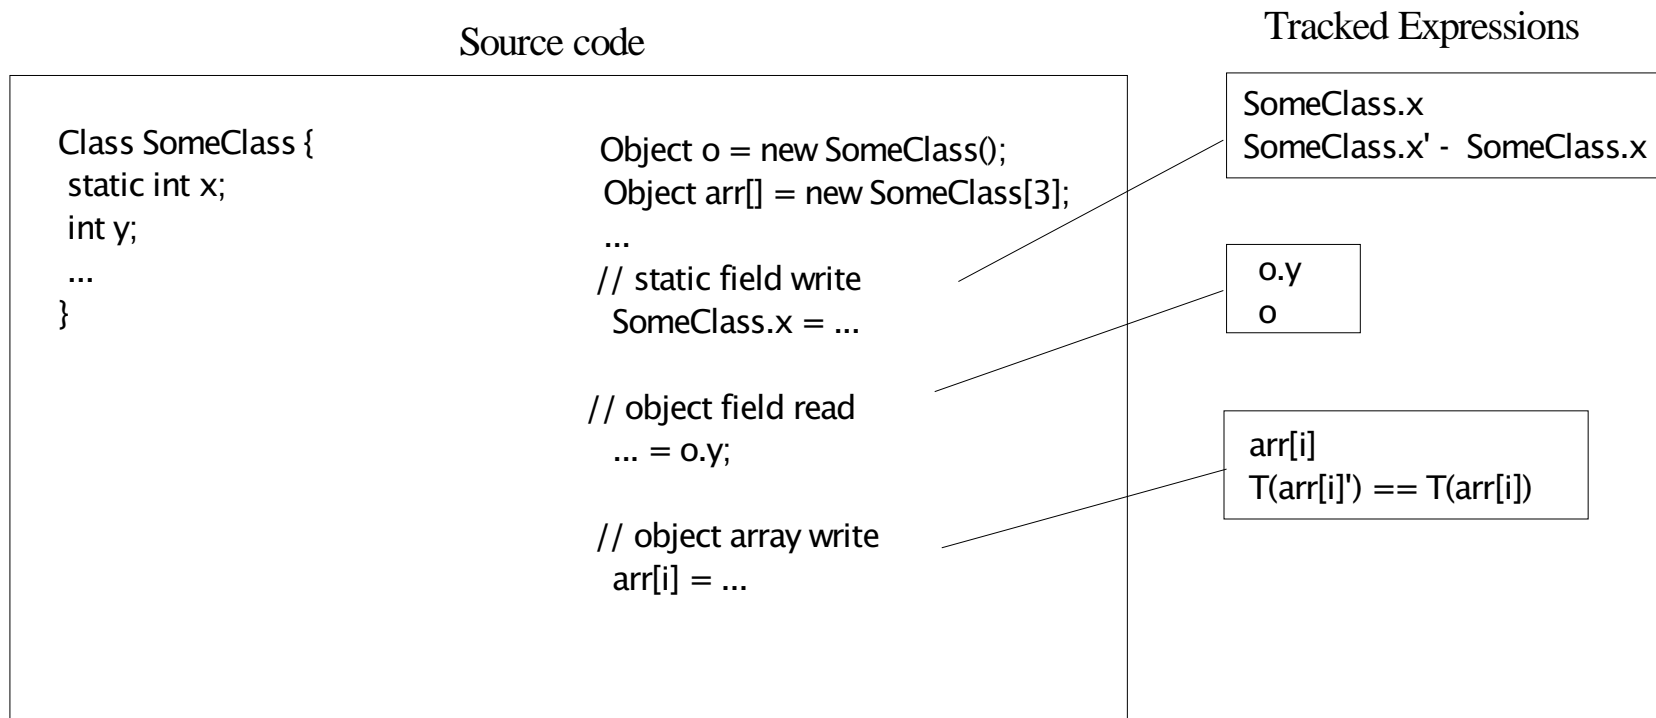
For writes, the variable name (e.g. o.x) refers to its value before the write, while the name with a ' suffix (e.g. o.x') refers to its value after the write.

Tracked Expressions



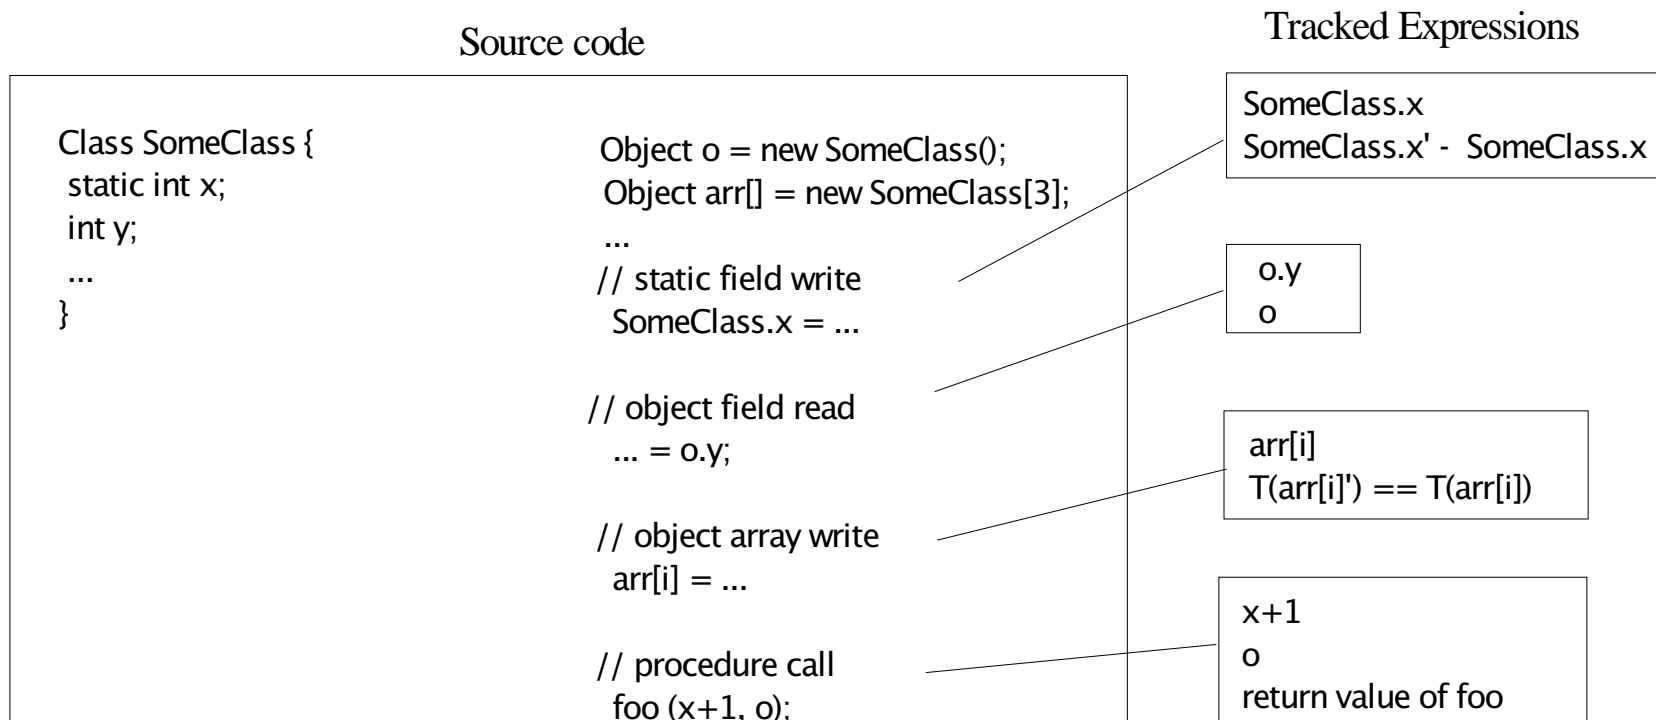
For writes, the variable name (e.g. o.x) refers to its value before the write, while the name with a ' suffix (e.g. o.x') refers to its value after the write.

Tracked Expressions



For writes, the variable name (e.g. o.x) refers to its value before the write, while the name with a ' suffix (e.g. o.x') refers to its value after the write.

Tracked Expressions



For writes, the variable name (e.g. o.x) refers to its value before the write, while the name with a ' suffix (e.g. o.x') refers to its value after the write.

Default Invariant Representation

- ◆ **Compact, for speed**
 - ◆ **Space overhead \propto static size of program**
- ◆ **For each tracked expression**
 - ◆ **Convert to integer**
 - **References map to hashcode of class name**
 - ◆ **Each expression stores sample value and mask**
 - **Mask tracks which bits have remained invariant**
 - ◆ **Meeting value incompatible with current hypothesis relaxes the mask/invariant**
 - **Mask can be relaxed up to #bits times**

Invariant Confidence

- ◆ **Defined as:**
 - ◆ **# Samples/# of bits marked invariant by mask**
- ◆ **High if same value observed many times, or with small number of bits change**
- ◆ **Users look at invariant violations with large confidence changes first**

User Extensions

- ❖ **Default modes work well in practice**
 - ❖ **our experiments run in this mode**
- ❖ **User can change defaults**
 - ❖ **based on class, method, target field/method name, read v/s write, type of value accessed, line number...**
 - ❖ **Change set of tracked expressions**
 - ❖ **Change invariant representation**
 - ❖ **Change confidence computation**
- ❖ **Writing a new invariant is 20-30 LoC**
 - ❖ **Examples: Min/Max values, Self loops**

DIDUCE Experiences

Program	Lines of code	# Classes instr.	# Program points instrumented	Slowdown
Sun MAJC MP Simulator	3300	10/28	3204	8–12 (10 proc)
MailManage + Javamail lib (SourceForge)	21700	214 /214	13014	6
JSSE lib (Java Secure Sockets Layer)	30000 +RSA libraries	384 /384	34844	8
Joeq JVM+JIT (SourceForge)	31500	18/137	3371	20

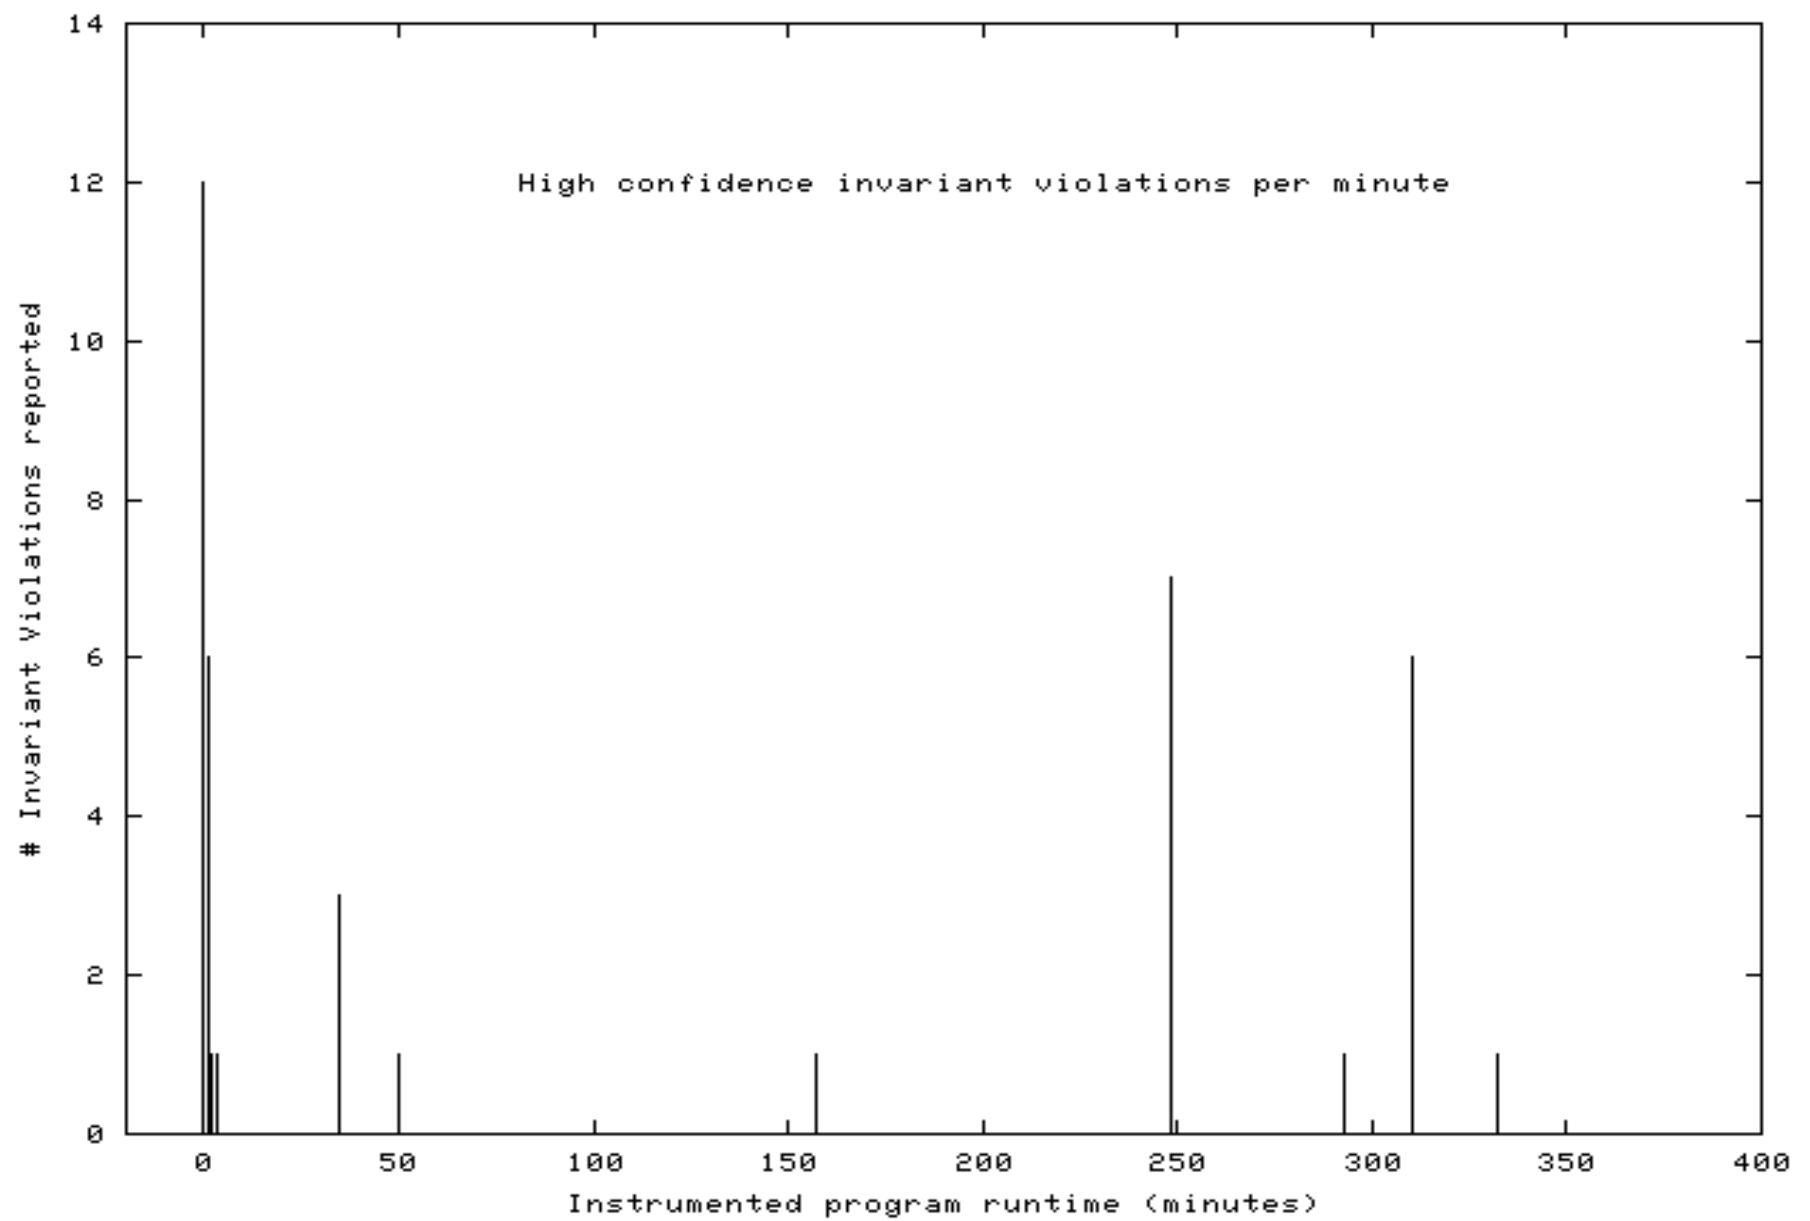
MAJC Simulator

- ◆ Simulator stable, in active use
- ◆ DIDUCE used initial part of run for training; ignored violations in this phase
- ◆ DIDUCE found 2 otherwise undetected errors
- ◆ DIDUCE accurately root-caused 3 errors
- ◆ Reported violation on 1 error was missed
 - ◆ Happened early, incorrect model was built!
- ◆ Reported 10 corner cases, all interesting to programmer!

Finding Simulator Errors

```
for (replaced = 0; replaced < assoc; replaced++) {  
    // Bug – shd have checked for 0 or 2  
    if (status[replaced][curset] == 0)  
        break;  
}
```

- ❖ **Occurred after 1 hour of uninstrumented execution time**
- ❖ **Another error detected even later: store to cache line in invalid state**



Debugging Known Errors

- ♦ **Correctly root-caused 3 errors which resulted in assertion failures**
- ♦ **One of these was extremely hard to understand for programmer**
 - ♦ **Occurred after an hour of execution time**
 - ♦ **Tried several iterations of debugging, gave up**
- ♦ **Ran DIDUCE overnight:**
 - ♦ **6 invariant violations reported just before error**
 - ♦ **Precisely pinpointed failing scenario**

JSSE Example

- ❖ **Java Secure Sockets Extension (JSSE) v. 1.0.2 code, ships with JDK 1.2/1.3**
- ❖ **Programmer tried adding a proxy server, changed timings**
- ❖ **Saw intermittent failures**
- ❖ **Spent 2 days working backwards from failure to root cause, through unfamiliar code**

JSSE Bug

```
InputStream s = x;  
// x is instance of SocketInputStream  
if (...) {  
    int len = ...; // const expr. = 74;  
    byte[] hdr = new byte[len];  
    s.read(hdr);  
}
```

- ❖ On passing runs, invariant on return value of `read()` as 74
- ❖ Failing run had a different value!
- ❖ Original programmer misunderstood contract with `InputStream.read(byte[])` - not guaranteed to completely fill array

Observations

- ♦ **Invariant violations tend to occur in clumps**
- ♦ **Many clues due to an anomaly**
 - ♦ **Can trade-off accuracy for overhead**
- ♦ **Debug forwards**
 - ♦ **Source of error may be far away from failure**
- ♦ **Debug backwards:**
 - ♦ **trail of interesting events along the way**
- ♦ **Random data tends to have low confidence**

Noise and Overhead

- ◆ **Users can select which program points to watch**
 - ◆ **Reduces instrumentation overhead**
 - ◆ **Also reduces noise due to unimportant invariants**
- ◆ **Overhead parallelized by using different machines with different parts instrumented**
- ◆ **No overhead on network or I/O operations**

Summary

- ◆ **Need automatic bug detection and debugging tools**
- ◆ **User assertions are insufficient**
- ◆ **DIDUCE approach finds hard bugs in large programs, real-life situations**
- ◆ **No up-front investment required**
- ◆ **Works for different application domains**
- ◆ **Knowing corner cases gives programmer better feel for what the program is doing**

Uses (1)

- ♦ **Automatic Root-Cause Analysis of bugs**
 - ♦ Check invariants learnt on runs which pass
- ♦ **Debug long-running programs**
 - ♦ Train on early part of program
- ♦ **Debug component-based software**
 - ♦ Train components on other configurations
- ♦ **Watch invariant violations when correct answer may be unknown**

Uses (2)

- ◆ **Understand legacy code**
 - ◆ "Can this ever happen ?"
- ◆ **Analyze test suite completeness**
 - ◆ Incorrect invariants indicate holes in test suite
- ◆ **Assist in program evolution**
 - ◆ "Did I break anything ?"

Questions ?

Contact Information:

sudheendra.hangal@sun.com

lam@cs.stanford.edu

DIDUCE now open sourced and available at:
<http://ditude.sf.net>

Backup: Mailmanage Example

- ♦ **Open source project to manipulate email**
 - ♦ **<http://sf.net/mailmanage>**
- ♦ **Uses JavaMail library**
- ♦ **Library threw cryptic IOException on 1 mailbox (of 300)**
- ♦ **Did not have source to JavaMail**
 - ♦ **Blindly instrumented all classes**
- ♦ **Trained on mailboxes which worked**
 - ♦ **Hit one invariant violation just before failure**

Backup: Mailmanage Error

```
do {  
    switch (buffer[index]) {  
        case 'E': ...; index +=...; break;  
        // other cases  
    } while (buffer[index++] != ')');
```

- ❖ **Obtained source code of class in question**
- ❖ **DIDUCE invariant violation was:**
 - ❖ **buffer[index] at this point = Ctrl-A (0x10), instead of ')' or ' '**
- ❖ **Bug: Length mismatch - Solaris IMAP server returning wrong data for a DOS attachment**

Backup: Mailmanage Lessons

- ❖ Error not in user code, not even in JavaMail library
- ❖ Error in IMAP server on a different machine
- ❖ Error detected as soon as it propagated into instrumentation domain
- ❖ DIDUCE helped zoom in to error in completely unfamiliar code!
- ❖ Helped to quickly identify faulty component and replace it
- ❖ Trade off accuracy for overhead

Backup: Joeq

- ♦ **Research Java VM, written in Java**
- ♦ **DIDUCE hit a bug, which caused an assert later:**
 - ♦ **JAR file had duplicate entries for same filename**
 - ♦ **Caused inconsistency between memory image and Jar file**
- ♦ **Detected because return value of Hashtable.put () not null!**

Backup: Future Work

- ◆ **Focus on strong invariants - enable invariant checking in deployed code ?**
 - ◆ **Overhead needs to be lowered**
 - ◆ **Identify critical "state" variables**
- ◆ **Get DIDUCE deployed in testing/bug analysis environments**